

1 What is CLAP?

CLAP = Contrastive Language–Audio Pretraining

👉 It is **the same idea as CLIP**, but instead of **image–text**, it uses **audio–text pairs**.

- CLIP → Image + Text
 - **CLAP → Audio + Text** ✓
-

2 Motivation (Why CLAP?)

Problem

- Audio understanding usually needs **manual labels**
- Labeling sounds is **expensive and limited**

Key Idea

Use **natural language descriptions** to teach the model what sounds mean

✓ Internet already has **millions of audio clips with captions or tags**

Examples:

- “engine knock sound”
- “bird chirping”
- “brake squeal”

These descriptions give **semantic meaning** to sounds.

3 Training Data Used

CLAP is trained on **audio–text datasets**, such as:

- **AudioCaps**
- **FreeSound**
- **WavCaps**

Each sample has:

- 🎧 Audio clip

- 🗒️ Text caption or tag
-

4 CLAP Architecture (Very Simple)

CLAP has **two encoders**, just like CLIP:

🎧 Audio Encoder

- CNN or Audio-ViT
- Converts audio → audio embedding (vector)

🗒️ Text Encoder

- Transformer
 - Converts text → text embedding (vector)
-

Architecture Flow (in words)

Audio Clip → Audio Encoder → Vector
Text → Text Encoder → Vector

Both vectors are mapped into the **same embedding space**.

5 Training Objective (Contrastive Loss)

How training works

- Batch has **N audio-text pairs**
- For one audio clip:
 - ✓ 1 correct text (positive)
 - ✗ N-1 wrong texts (negatives)
- Similarly for text → audio

👉 Total negatives $\approx N^2 - N$

Goal of Contrastive Loss

- **Increase similarity** between matching audio–text pairs
- **Decrease similarity** between mismatched pairs

✦ Same contrastive idea as NT-Xent / CLIP.

6 What Can CLAP Do After Training?

✓ Zero-Shot Audio Classification

- No retraining
- Just compare audio with text prompts

Example prompts:

- “engine sound”
- “rain noise”
- “brake squeal”

Pick the text with **highest similarity**.

✓ Audio–Text Retrieval

- Given audio → find best matching text
 - Given text → find relevant audio clips
-

✓ Sound Event Detection

- Detect sound types using **text queries**
 - Example: “car horn”, “glass breaking”
-

7 Key Paper

■ LAION-CLAP (2023)

- Large-scale audio–language contrastive pretraining
 - Similar impact to CLIP but for audio
-

8 Real-World / AV Applications

CLAP is useful in **autonomous vehicles (AV)**:

- 🚗 Classify **road surface sounds**
- 🚗 Detect **brake squeal**
- 🔧 Identify **engine anomalies** using text queries

Example query:

"engine knocking sound"

CLAP finds matching audio patterns **without retraining**.

9 One-Line Intuition (Very Important!)

CLAP learns to match sounds with their language descriptions, enabling zero-shot audio understanding.

10 One-Paragraph Exam Answer ✓

CLAP is a contrastive learning model that jointly trains an audio encoder and a text encoder using large-scale audio–text pairs. By maximizing similarity between matching audio and text embeddings and minimizing it for mismatched pairs, CLAP learns a shared semantic space. This enables zero-shot audio classification, audio–text retrieval, and sound event detection without task-specific training.